

# Pattern Recognition

## Assignment 2

---

### Team Members

Name	ID
Ramy Ahmed ElSayed	19015649
Zyad Samy Ramadan	19015720
Abdelmoniem Hany	19017359

### Problem Statement

The exponential growth of network traffic has led to an increase in network anomalies, such as cyber attacks, network failures, and hardware malfunctions. Network anomaly detection is a critical task for maintaining the security and stability of computer networks. The objective of this assignment is to help students understand how K-Means and Normalized Cut algorithms can be used for network anomaly detection.

### Data Set

The dataset used in the assignment is from the "KDD Cup 1999" dataset and is partitioned into the following:

- 10% Training Data and its labels.
- Correction Testing Data and its labels.
- Full Data for Spectral Clustering and its labels.

```
def loadData():
    trainingData = pd.read_csv('drive/MyDrive/Anomaly Detection
Data/kddcup.data_10_percent', sep=",", header = None)
    trainingDataSpectral = pd.read_csv('drive/MyDrive/Anomaly Detection
Data/kddcup.data/kddcup.data', sep=",", header = None)
    testingData = pd.read_csv('drive/MyDrive/Anomaly Detection Data/corrected/corrected',
sep=",", header = None)
    stringCols = [1, 2, 3, 41]
    for i in stringCols:
```

```

trainingData[i] = pd.factorize(trainingData[i])[0]
testingData[i] = pd.factorize(testingData[i])[0]
trainingDataSpectral[i] = pd.factorize(trainingDataSpectral[i])[0]
return trainingData.loc[:,0:40], trainingData.loc[:, 41:], testingData.loc[:,0:40],
testingData.loc[:, 41:], trainingDataSpectral

```

## K-Means

The implementation of K-means is as follows:

1. Select a random subset of the points to be the initial centroids  
*by generating a permutation of the numbers from 1:n\_rows and then selecting those points from the dataset as the centroids.*
2. Iterate until
  - a. Convergence: when the delta (change in centroids) becomes less than some given threshold epsilon.
  - b. Max iterations are exceeded.
3. During each iteration:
  - a. Calculate the proximity matrix (n,k) which tells how far each point is from each centroid.
  - b. Select the cluster for each point based on the closest centroid.
  - c. Update each cluster centroid.
  - d. Calculate the change in centroids to detect convergence.

```

def kMeans(data, k, epsilon, iterations):
    n_rows = data.shape[0]
    random_idx = np.random.RandomState(42).permutation(n_rows)
    centroids = data[random_idx[:k]]
    clusters = np.zeros(n_rows)
    proximityMatrix = np.zeros((n_rows, k))
    delta = np.inf
    iteration = 0

    while delta > epsilon and iteration < iterations:
        for i in range(k):
            proximityMatrix[:,i] = np.linalg.norm(data - centroids[i], axis=1)

        clusters = np.argmin(proximityMatrix, axis = 1)

        iteration += 1
        old_centroids = deepcopy(centroids)

        for i in range(k):

```

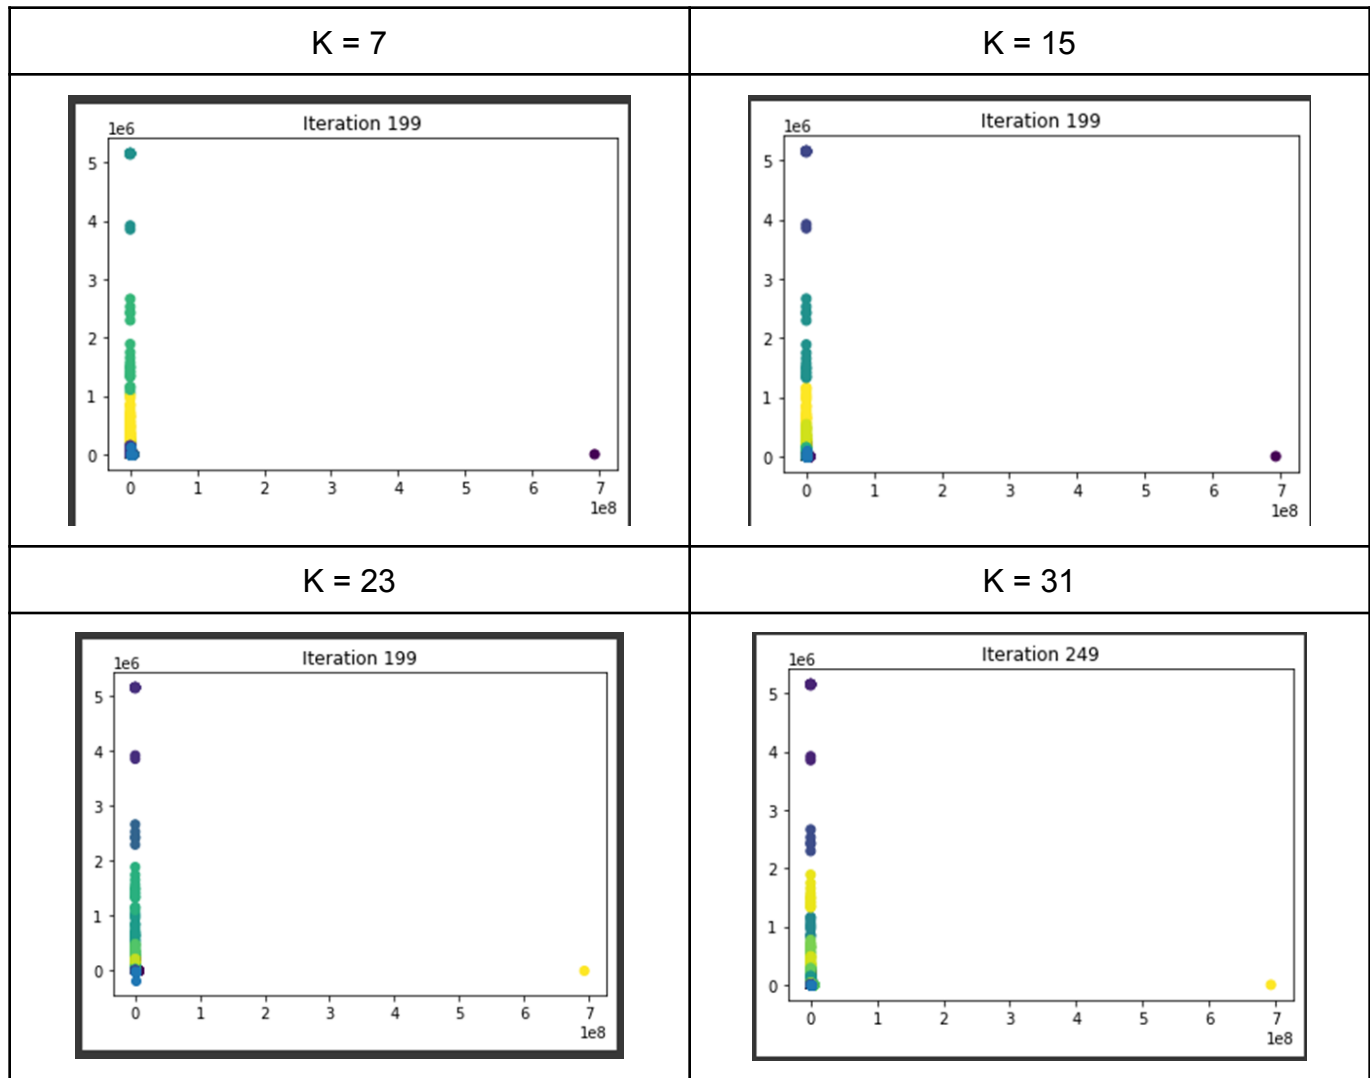
```

centroids[i,:] = np.mean(data[clusters == i,:], axis = 0)
if np.isnan(centroids).any():
    centroids[i] = old_centroids[i]

delta = np.linalg.norm(centroids - old_centroids)
return clusters, centroids

```

## Clusters Visualized



## K-Means Evaluation

The model was trained using the 10% dataset “*kddcup.data\_10\_percent*” on different values of  $K = [7, 15, 23, 31, 45]$  and was then tested on the “*corrected.gz*” dataset and it produced the following measures:

```
actual_labels = testingLabels.to_numpy().flatten()
for k in kArray:
    clusters, centroids = kMeans(trainingData.to_numpy(), k, 0.001, 200)
    predicted_labels = predict(centroids, testingData)
    print(f'K-Means with k={k}')
    checkClustering(predicted_labels, actual_labels, k)
    print()
```

```
K-Means with k=7
Precision Score: 0.35724224700025636
Recall Score: 0.9994639263064989
F Score: 0.5263493980559961
Conditional Entropy: 0.7512333969826389

K-Means with k=15
Precision Score: 0.6311021489567514
Recall Score: 0.6146898689806143
F Score: 0.6227878998581184
Conditional Entropy: 0.3720591818015853

K-Means with k=23
Precision Score: 0.6227486330572504
Recall Score: 0.4462170403306566
F Score: 0.5199064082798301
Conditional Entropy: 0.3981877037228226

K-Means with k=31
Precision Score: 0.7286580436816733
Recall Score: 0.42186554122561254
F Score: 0.5343579636239033
Conditional Entropy: 0.42213098103853636

K-Means with k=45
Precision Score: 0.9218192366460805
Recall Score: 0.36134602791237447
F Score: 0.5191782053575126
Conditional Entropy: 0.24592909396855236
```

## Spectral Clustering using K-ways Normalized Cut

```
def SpectralClustering(data, clustersNumber, gammaValue = 1):
    similarityMatrix = rbf(data, gamma=gammaValue)
    degreeMatrix = np.diag(np.sum(similarityMatrix, axis=1))
    laplacianMatrix = degreeMatrix - similarityMatrix
    normAsymLaplacianMatrix = np.dot(sc.linalg.inv(degreeMatrix), laplacianMatrix)
    eigenValues, eigenVectors = sc.linalg.eig(normAsymLaplacianMatrix)
    eigenValues = np.abs(np.real(eigenValues))
    eigenVectors = np.real(eigenVectors)
    sort_perm = eigenValues.argsort()
    eigenValues = eigenValues[sort_perm]
    eigenVectors = eigenVectors[:, sort_perm]
    eigenValues = eigenValues[:clustersNumber]
    U = eigenVectors[:, :clustersNumber]
    Y = normalize(U, axis=1)
    labels = KMeans(n_clusters=clustersNumber, n_init="auto",
random_state=42).fit_predict(Y)
    return eigenValues, Y, labels
```

In order to compare spectral clustering with k-means we will run both on k=23, and on a smaller subset of the dataset (0.15%) so that it can run on the available resources.

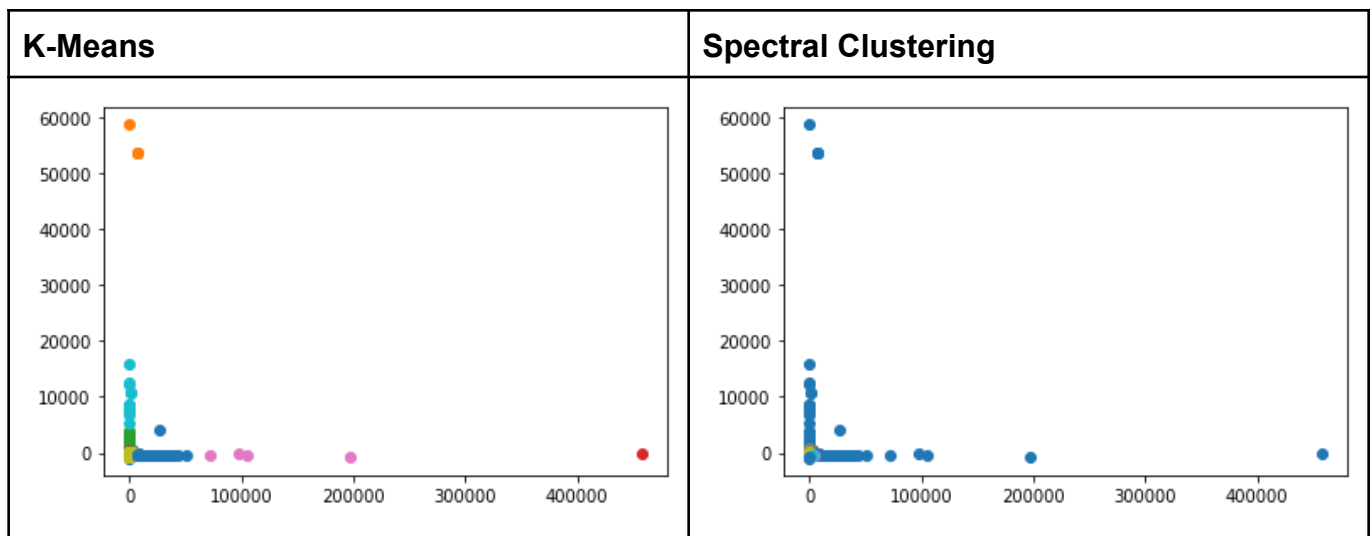
```
X_train_spectral, _ = train_test_split( totalDataSpectral, train_size=0.0015,
random_state=42)
X_train_meanShift, _ = train_test_split( totalDataSpectral, train_size=0.00015,
random_state=42)

spectral_train = X_train_spectral.loc[:,0:40]
spectral_labels = X_train_spectral.loc[:, 41:]
```

And the results was as follows:

	K-Means	Spectral Clustering
Precision	0.8034	0.4210
Recall	0.7236	0.9987
F Score	0.7614	0.5923
Conditional Entropy	0.2934	0.9945

And then PCA was used on the data to apply dimensionality reduction to visualize the results



## New Clustering Algorithm (Mean Shift)

```
def mean_shift(data, bandwidth = None, bandwidthStep = 100):

    if bandwidth is None:
        all_data_centroid = np.average(data,axis=0)
        all_data_norm = np.linalg.norm(all_data_centroid)
        bandwidth = all_data_norm/bandwidthStep

    centroids = {}

    for i in range(len(data)):
        centroids[i] = data[i]

    weights = [i for i in range(bandwidthStep)][::-1]
    while True:
        new_centroids = []
        for i in centroids:
            in_bandwidth = []
            centroid = centroids[i]

            for featureset in data:

                distance = np.linalg.norm(featureset-centroid)
                if distance == 0:
                    distance = 0.000000000001
                weight_index = int(distance/bandwidth)
                if weight_index > bandwidthStep - 1:
                    weight_index = bandwidthStep - 1

                to_add = (weights[weight_index]**2)*[featureset]
                in_bandwidth +=to_add

            new_centroid = np.average(in_bandwidth,axis=0)
            new_centroids.append(tuple(new_centroid))

        uniques = sorted(list(set(new_centroids)))

        to_pop = []

        for i in uniques:
            for j in [i for i in uniques]:
                if i == j:
                    pass
                elif np.linalg.norm(np.array(i)-np.array(j)) <= bandwidth:
                    to_pop.append(j)
                    break
```

```

for i in to_pop:
    try:
        uniques.remove(i)
    except:
        pass

prev_centroids = dict(centroids)
centroids = {}
for i in range(len(uniques)):
    centroids[i] = np.array(uniques[i])

optimized = True

for i in centroids:
    if not np.array_equal(centroids[i], prev_centroids[i]):
        optimized = False

if optimized:
    break

proximityMatrix = np.zeros((data.shape[0], len(centroids)))
for i in range(len(centroids)):
    proximityMatrix[:,i] = np.linalg.norm(data - centroids[i], axis=1)

clusters = np.argmin(proximityMatrix, axis = 1)

return clusters

```

## Mean Shift Evaluation

```

Precision Score:  0.47576078664067134
Recall Score:    0.9584108877922166
F Score:        0.6358713200514358
Conditional Entropy:  0.8557752609213186

```



## Evaluation

The following is the code used for the evaluation of the models

```
from sklearn.metrics.cluster import pair_confusion_matrix

def checkClustering(resultingLabels, trueLabels, K):
    confusionMat = pair_confusion_matrix(trueLabels, resultingLabels)
    precision = confusionMat[1][1] / (confusionMat[1][1] + confusionMat[0][1])
    recall = confusionMat[1][1] / (confusionMat[1][1] + confusionMat[1][0])
    f1score = 2 * precision * recall / (precision + recall)
    print("Precision Score: " , precision)
    print("Recall Score: " , recall)
    print("F Score: " , f1score)
    print("Conditional Entropy: " , conditionalEntropy(resultingLabels, trueLabels, K))
```

```
def conditionalEntropy(predictedLabels, trueLabels, K):
    predictedLabeledClusters = [[] for _ in range(K)]
    for i in range(len(predictedLabels)):
        predictedLabeledClusters[predictedLabels[i]].append(trueLabels[i])

    entropy = np.zeros(K)
    N = 0

    for i in range(len(predictedLabeledClusters)):
        N += len(predictedLabeledClusters[i])
        count = np.zeros(K)
        for j in range(len(predictedLabeledClusters[i])):
            for k in range(len(predictedLabeledClusters[i][j])):
                if j == predictedLabeledClusters[i][j][k]:
                    count[j] += 1
        for j in range(len(predictedLabeledClusters[i])):
            if count[j] != 0:
                entropy[i] += (- count[j]/len(predictedLabeledClusters[i])) *
math.log(count[j]/len(predictedLabeledClusters[i]))

    totalEntropy = 0
    for i in range(len(predictedLabeledClusters)):
        totalEntropy += (len(predictedLabeledClusters[i]) / N) * entropy[i]

    return totalEntropy
```